

Effect of Key Match Events on Football Passmaps

Vid Stropnik^{a,1}, Vuk Đuranović^a, and Maruša Oražem^a

^aUniversity of Ljubljana, Faculty of Computer and Information Science, Večna pot 113, SI-1000 Ljubljana, Slovenia

The manuscript was compiled on June 21, 2021

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Association football is the world's most popular sport. During gameplay, players attempt to create goal-scoring opportunities through various methods of individual control of the ball. While these can include dribbling, tackling and taking direct shots, the most common one is passing the ball to a teammate. Progressing the ball along the pitch with a series of passes helps the team retain control of the match, enabling them to realize set plays. It is worth noting that this method is not only common, but also effective, as exhibited by the correlation of per-game average pass frequency and overall tea, success. (?)

Ball movement and goal-scoring opportunities in a football match are complex phenomena to model. Pass characteristics, such as their length, speed or the involved players' positions are often used to provide further insights into their importance and quality. Logging such characteristics for hundreds of passes that happen in any given football match, however, renders the resulting data structure very complex and hard to interpret. To infer more general information about longer stints of play, simpler arrangements, such as networks with player nodes and directed pass links, should be considered.

There are several events in any given football match, that might incur a fundamental change in a team's tactics and approach towards the game. The coaching staff might recognize weak points in the opposition's strategy and relay them to the players during half-time. Suddenly conceding a goal in a must-win game might prompt a team to take more risks and try a different approach, while we can often observe leading teams resort to time-wasting tactics to ensure their victory.

In this course project, we introduce novel ways of football match analysis using established network science approaches. By using methods of measuring change in team pass map centrality, by observing alterations of its underlying communities and by recognizing the (dis)appearance of distinct player graphlets as a result of a key event in a football match, we study common team behaviours that key match events might induce. We recognize these insights as useful in several areas. First, they serve as a form of establishing several ground

truths about common trends of play, that will underpin further research. Secondly, more fine-grained difference of successful and unsuccessful responses to key events might be observed and be used in match preparations by domain experts. Finally, we recognize real-time applications of the proposed methods as useful in broadcasting, betting, commentary and analytics, to name only a few fields.

This report is structured as follows: In the following paragraphs, we provide a brief overview of the work, relating to inferring insights from sports data using network analysis. Here, we also introduce some key academic contributions to our work in the tree network science fields of special interest: measuring centrality, detecting communities and gathering semantic information from network graphlets. A subsection corresponds to each of these fields in the succeeding *Results* section. In them, we introduce the main revelations about the effects of key match events on passing trends. We later comment on the implications of these results and asses their appropriateness for further application, while finally introducing our detailed methodologies at the end of the paper.

Related work

Evaluations of group performances based on their underlying network structures have been conducted for some time. Researchers were keen on understanding what kind of social, interaction or communication network structure boosts the productivity of a group in a working environment. Bavelas (?) was one of the pioneers in exploration of information transfer between groups. He discovered that the rate of information diffusion was significantly smaller in decentralized networks in comparison to centralized ones. It was later discovered by Shaw (?) that this gave an edge to decentralized groups over the centralized ones in solving complex tasks.

Team performance in collective sports has become a well researched area in the network analysis community in recent years. We recognize the conclusions by Grund (?) as a good starting point for the study of this field. In his work, he suggests that small networks (i.e. the eleven-node football team) are better explained with the use of the weighted edges between players. In such cases, weights are determined by the frequency of passes between two players. He also introduces a measure called network intensity, as a counterpart to network density in large and unweighted networks. This measure is computed as a function of weights for each link between players, nicely called the passing-rate of the team. As a result of his work, Grund confirmed that an increase in team passing rate leads to increased performance, while increases in centralization of play reduces team performance. Similarly to our work, network metrics were used to examine differences

All authors contributed equally to this work.

¹To whom correspondence should be addressed. E-mail: vs2658@student.uni-lj.si

between halftimes in work by Clementine Clementine, albeit with a different form of pass-map representation. The authors of (?), again envisioned a football match as a weighted directed network, where they introduced a node measure called flow centrality - a function of pass accuracy and ratio of times a player was included in an action, leading to a shot on goal. Using it, they successfully graded individual player performances within a team. The measure is based on betweenness centrality, which has also been applied to football matches to examine correlations of physical demands in elite football players (?). While these works' conclusions are introduced as solid and objective, the importance of considering additional meta-data in the inference of flow using centrality measures was emphasized in the traffic domain in (?).

While several graphlet and motif-related methodology in our work is heavily inspired by the work of Pržulj (?), the concept of network motifs, initially shown by Milo et al. (?), has been successfully applied to sporting pass-maps by other authors as well. Notably, Penna and Touchette (?) analyzed the network motifs to analyze passing behaviour, while Bekkers and Shaunak (?) use meta-data equipped motifs (Possession Motifs and Expected Goal Motifs) to answer more in depth domain questions and graphically visualize team or player similarity by means of radar graph.

Due to advances in player-tracking technology, most recent research, such as those by Fernandez and Born (?), or by Chawla et al. (?), exploits the spatial-temporal data about observed passes. Due to time restrictions, posed by the gathered data, we don't consider these additional pieces of information in our analysis, but instead focus on the starting formation of the team in determining the positions of players. Researchers opting for the former approach of inferring information from additional metadata, however, should also consider work by Mattson and Takes (?) on pass trajectory extraction. Spatio-temporal methods are also successfully applied to football to accurately detect individual player behaviour and changes in team formation. (?).

Results

Centrality and intensity.

Communities. One of our assumptions was that the number of communities one might observe would change after a significant event in soccer game such as goal event, or the event of a player receiving a red card. On a small weighted graph, such a football squad is, it is hard to observe drastically changes in number of communities, as their number is limited by the number of players that can be on the field in the same time. After testing all key events in order to check whether number of communities change, we can state that we found no significant results.

In contrast to that, we were able to find very interesting results regarding the structure of communities that occur in a game. We uncovered the proportions of players by position that appear in communities of different sizes.

Network fragments.




Fig. 1. Stacked barplot showing proportions of players by their position that appear in different communities. From this it can be observed that attackers tend to appear in communities of smaller size.




Fig. 2. The distribution of graph edge weights in Bayern Munich's 2012 Champions league final loss versus Chelsea shows that edge weights roughly follow a Poisson distribution. As the same trends emerge in other examined games as well, we note that pruning the graphs at just above the median $\lfloor \lambda + \frac{1}{3} - \frac{0.02}{\lambda} \rfloor \approx \lambda$ will result in considering only the non-trivial passes, carrying information about frequent passes of play.

Discussion

Methods

All of the described results were achieved by analysing directed networks with player-nodes and weighted pass edges, with weights corresponding to pass frequency.

Data. All of the examined networks were parsed from alternative data struces by Statsbomb (?). Unless noted otherwise, the described results were achieved by aggregating the findings of analyses, conducted on 789 parsed matches. For each successful pass, a directed edge was generated between the players, logged in the corresponding pass attempt and reception events, logged by the data provider. For the three considered key events - the half time break, the first goal and the first dismissal in any given match, 6 directed *pajek* files were created for each involved team (3 before and 3 after the key event), where games without certain events contain empty *after-event* files. Additionally, the graphs' flow centrality variants were created and saved according to the procedure, described in the following subsection. For the calculation of network intensity, the description of which also follows, the team possession time was calculated as the sum of timestamp differences in each uninterrupted possession stint.

Centrality and intensity.

Communities. For finding communities in weighted graph we used Leiden community detection algorithm. Prior to performing the search, we pruned the graphs, so that edges having a value of their weight smaller than the median of all values of weights in a graph.

In order to check for any valuable information we may have missed with our algorithm selection, we ran additional experiments using CPM algorithm.

Determining how do communities detected in certain periods of a game differ from standard playing lines required from us to manually input correct formations (line-ups) of teams in observed matches. Since this is a very time-consuming task, we restricted ourselves to 20 matches for which we described the formation in which teams played. Using the parsing program which we developed earlier, we constructed weighted graphs for each of the teams, two graphs representing the state of the game prior and past a certain key event for each team, four in total for a given event in a football match.

Similarity between the detected communities and communities pre-defined by position was measured using NMI (normalized mutual information). Since most of the time, graphs would have NMI around 0.3, we can only conclude that communities generated by passmaps between players differ from the artificial communities imposed by coach's formation.

Network fragments. In all experiments concerning network fragments, the networks were pruned before examination. This was done by removing all edges with weights equal to or less than the median weight from the observed graph. The estimation of an adequate pruning point stems from the insights, described in the caption of Figure ??.

To examine the orbit-representation in the observed graphs, an atlas of all possible directed graphlets of node-count 3 or less was created. From there, each unique edge subset from the analysed graph was extracted and a hash table of benchmark orbital roles created. For each subgraph, inferred from the edge subsets, the orbital roles for each contributing node were calculated. The resulting table of this procedure denoted the number of occurrences of each player in each of the 34 considered orbital roles. Each network's orbital profile is then calculated as the average number of player-occurrences in each orbital role, divided by 11 - the number of players on the pitch. By comparing the differences between these profiles before and after key events, the final insights of the distribution of PPO (Player per orbit) differences is achieved.

ACKNOWLEDGMENTS. The authors would like to acknowledge the help of Doc. Dr. Lovro Šubelj and Assoc. Prof. Dr. Erik Štrumbelj in the identification of feasible problems within our domain of interest for this course project and for all the help offered in solving them.